

```
In [1]: import os
import re
import csv
import sys
import codecs
import numpy as np
import pandas as pd
import keras.layers as KL
from nltk.corpus import stopwords
from nltk.stem import SnowballStemmer
from string import punctuation
from gensim.models import KeyedVectors
from keras import backend as KB
from keras.preprocessing.text import Tokenizer
from keras.preprocessing.sequence import pad_sequences
from keras.layers import Dense, Input, LSTM, Embedding, Dropout, Activation
from keras.layers.merge import concatenate
from keras.models import Model
from keras.layers.core import Reshape, Permute, Lambda, RepeatVector
from keras.layers.normalization import BatchNormalization
from keras.callbacks import EarlyStopping, ModelCheckpoint
```

```
C:\Users\tianh\Desktop\environments\mlenv\lib\site-packages\h5py\__init__.py:
36: FutureWarning: Conversion of the second argument of issubdtype from `float` to `np.floating` is deprecated. In future, it will be treated as `np.float64 == np.dtype(float).type`.
    from ._conv import register_converters as _register_converters
Using TensorFlow backend.
C:\Users\tianh\Desktop\environments\mlenv\lib\site-packages\gensim\utils.py:1
167: UserWarning: detected Windows; aliasing chunkize to chunkize_serial
    warnings.warn("detected Windows; aliasing chunkize to chunkize_serial")
```

```
In [2]: BASE_DIR = 'data/'
EMBEDDING_FILE = BASE_DIR + 'GoogleNews-vectors-negative300.bin'
TRAIN_DATA_FILE = BASE_DIR + 'train.csv'
TEST_DATA_FILE = BASE_DIR + 'test.csv'
MAX_SEQUENCE_LENGTH = 30
MAX_NB_WORDS = 100000
EMBEDDING_DIM = 300
VALIDATION_SPLIT = 0.3
LOAD_DATA = False
```

```
In [3]: num_lstm = np.random.randint(175, 275)
num_dense = np.random.randint(100, 150)
rate_drop_lstm = 0.15 + np.random.rand() * 0.25
rate_drop_dense = 0.15 + np.random.rand() * 0.25
act = 'relu'
re_weight = True # whether to re-weight classes to fit the 17.5% share in test set
STAMP = 'lstm'
```

```
In [4]: print('Indexing word vectors')
word2vec = KeyedVectors.load_word2vec_format(EMBEDDING_FILE, binary=True)
print('Found %s word vectors of word2vec' % len(word2vec.vocab))
```

```
Indexing word vectors
Found 3000000 word vectors of word2vec
```

```
In [5]: print('Processing text dataset')

# The function "text_to_wordlist" is from
# https://www.kaggle.com/currie32/quora-question-pairs/the-importance-of-cl
eaning-text
def text_to_wordlist(text, remove_stopwords=False, stem_words=False):
    # Clean the text, with the option to remove stopwords and to stem words.

    # Convert words to lower case and split them
    text = text.lower().split()

    # Optionally, remove stop words
    if remove_stopwords:
        stops = set(stopwords.words("english"))
        text = [w for w in text if not w in stops]

    text = " ".join(text)

    # Clean the text
    text = re.sub(r"[^A-Za-z0-9^,!./'+=]", " ", text)
    text = re.sub(r"what's", "what is ", text)
    text = re.sub(r"\'s", " ", text)
    text = re.sub(r"\'ve", " have ", text)
    text = re.sub(r"can't", "cannot ", text)
    text = re.sub(r"\n't", " not ", text)
    text = re.sub(r"i'm", "i am ", text)
    text = re.sub(r"\re", " are ", text)
    text = re.sub(r"\d", " would ", text)
    text = re.sub(r"\ll", " will ", text)
    text = re.sub(r",", " ", text)
    text = re.sub(r"\.", " ", text)
    text = re.sub(r"!", " ! ", text)
    text = re.sub(r"\/", " ", text)
    text = re.sub(r"\^", " ^ ", text)
    text = re.sub(r"\+", " + ", text)
    text = re.sub(r"\-", " - ", text)
    text = re.sub(r"\=", " = ", text)
    text = re.sub(r"\'", " ", text)
    text = re.sub(r"(\d+)(k)", r"\g<1>000", text)
    text = re.sub(r":", " : ", text)
    text = re.sub(r" e g ", " eg ", text)
    text = re.sub(r" b g ", " bg ", text)
    text = re.sub(r" u s ", " american ", text)
    text = re.sub(r"\0s", "0", text)
    text = re.sub(r" 9 11 ", "911", text)
    text = re.sub(r"e - mail", "email", text)
    text = re.sub(r"j k", "jk", text)
    text = re.sub(r"\s{2,}", " ", text)
```

```
# Optionally, shorten words to their stems
if stem_words:
    text = text.split()
    stemmer = SnowballStemmer('english')
    stemmed_words = [stemmer.stem(word) for word in text]
    text = " ".join(stemmed_words)

# Return a list of words
return(text)

texts_1 = []
texts_2 = []
labels = []
with codecs.open(TRAIN_DATA_FILE, encoding='utf-8') as f:
    reader = csv.reader(f, delimiter=',')
    header = next(reader)
    for values in reader:
        texts_1.append(text_to_wordlist(values[3]))
        texts_2.append(text_to_wordlist(values[4]))
        labels.append(int(values[5]))
print('Found %s texts in train.csv' % len(texts_1))

tokenizer = Tokenizer(num_words=MAX_NB_WORDS)
tokenizer.fit_on_texts(texts_1 + texts_2)

sequences_1 = tokenizer.texts_to_sequences(texts_1)
sequences_2 = tokenizer.texts_to_sequences(texts_2)

word_index = tokenizer.word_index
print('Found %s unique tokens' % len(word_index))

data_1 = pad_sequences(sequences_1, maxlen=MAX_SEQUENCE_LENGTH)
data_2 = pad_sequences(sequences_2, maxlen=MAX_SEQUENCE_LENGTH)
labels = np.array(labels)
print('Shape of data tensor:', data_1.shape)
print('Shape of label tensor:', labels.shape)
```

```
Processing text dataset
Found 404290 texts in train.csv
Found 85518 unique tokens
Shape of data tensor: (404290, 30)
Shape of label tensor: (404290,)
```

```
In [6]: print('Preparing embedding matrix')

nb_words = min(MAX_NB_WORDS, len(word_index))+1

embedding_matrix = np.zeros((nb_words, EMBEDDING_DIM))
for word, i in word_index.items():
    if word in word2vec.vocab:
        embedding_matrix[i] = word2vec.word_vec(word)
print('Null word embeddings: %d' % np.sum(np.sum(embedding_matrix, axis=1) == 0))
```

```
Preparing embedding matrix
Null word embeddings: 37391
```

```
In [7]: perm = np.random.permutation(len(data_1))
idx_train = perm[:int(len(data_1)*(1-VALIDATION_SPLIT))]
idx_val = perm[int(len(data_1)*(1-VALIDATION_SPLIT)):]

data_1_train = np.vstack((data_1[idx_train], data_2[idx_train]))
data_2_train = np.vstack((data_2[idx_train], data_1[idx_train]))
labels_train = np.concatenate((labels[idx_train], labels[idx_train]))

data_1_val = np.vstack((data_1[idx_val], data_2[idx_val]))
data_2_val = np.vstack((data_2[idx_val], data_1[idx_val]))
labels_val = np.concatenate((labels[idx_val], labels[idx_val]))

weight_val = np.ones(len(labels_val))
if re_weight:
    weight_val *= 0.472001959
    weight_val[labels_val==0] = 1.309028344
```

```
In [8]: embedding_layer = Embedding(nb_words,
    EMBEDDING_DIM,
    weights=[embedding_matrix],
    input_length=MAX_SEQUENCE_LENGTH,
    trainable=False)
lstm_layer = LSTM(num_lstm, dropout=rate_drop_lstm, recurrent_dropout=rate_dro
p_lstm)

sequence_1_input = Input(shape=(MAX_SEQUENCE_LENGTH,), dtype='int32')
embedded_sequences_1 = embedding_layer(sequence_1_input)
x1 = lstm_layer(embedded_sequences_1)

sequence_2_input = Input(shape=(MAX_SEQUENCE_LENGTH,), dtype='int32')
embedded_sequences_2 = embedding_layer(sequence_2_input)
y1 = lstm_layer(embedded_sequences_2)

merged = concatenate([x1, y1])
merged = Dropout(rate_drop_dense)(merged)
merged = BatchNormalization()(merged)

merged = Dense(num_dense, activation=act)(merged)
merged = Dropout(rate_drop_dense)(merged)
merged = BatchNormalization()(merged)

preds = Dense(1, activation='sigmoid')(merged)
```

```
In [9]: if re_weight:
    class_weight = {0: 1.309028344, 1: 0.472001959}
else:
    class_weight = None
```

```
In [10]: model = Model(inputs=[sequence_1_input, sequence_2_input], \
                     outputs=preds)
model.compile(loss='binary_crossentropy',
              optimizer='adam',
              metrics=['acc'])
model.summary()
print(STAMP)
if LOAD_DATA:
    model.load_weights(bst_model_path)
early_stopping = EarlyStopping(monitor='val_loss', patience=5)
bst_model_path = STAMP + '.h5'
model_checkpoint = ModelCheckpoint(bst_model_path, save_best_only=True, save_weights_only=True)

hist = model.fit([data_1_train, data_2_train], labels_train, \
                  validation_data=([data_1_val, data_2_val], labels_val, weight_val),
                  \
                  epochs=50, batch_size=2048, shuffle=True,
                  class_weight=class_weight, callbacks=[model_checkpoint])

model.load_weights(bst_model_path)
bst_val_score = min(hist.history['val_loss'])
```

Layer (type)	Output Shape	Param #	Connected to
input_1 (InputLayer)	(None, 30)	0	
input_2 (InputLayer)	(None, 30)	0	
embedding_1 (Embedding)	(None, 30, 300)	25655700	input_1[0] input_2[0]
lstm_1 (LSTM)	(None, 265)	599960	embedding_1[0][0] embedding_1[1][0]
concatenate_1 (Concatenate)	(None, 530)	0	lstm_1[0] lstm_1[1]
dropout_1 (Dropout)	(None, 530)	0	concatenate_1[0][0]
batch_normalization_1 (BatchNormalization)	(None, 530)	2120	dropout_1
dense_1 (Dense)	(None, 145)	76995	batch_normalization_1[0][0]
dropout_2 (Dropout)	(None, 145)	0	dense_1[0]
batch_normalization_2 (BatchNormalization)	(None, 145)	580	dropout_2
dense_2 (Dense)	(None, 1)	146	batch_normalization_2[0][0]
Total params:	26,335,501		

Trainable params: 678,451
Non-trainable params: 25,657,050

lstm
Train on 566006 samples, validate on 242574 samples
Epoch 1/50
566006/566006 [=====] - 124s 218us/step - loss: 0.
4395 - acc: 0.6688 - val_loss: 0.3643 - val_acc: 0.6976
Epoch 2/50
566006/566006 [=====] - 121s 215us/step - loss: 0.
3726 - acc: 0.6956 - val_loss: 0.3481 - val_acc: 0.7102
Epoch 3/50
566006/566006 [=====] - 122s 216us/step - loss: 0.
3516 - acc: 0.7155 - val_loss: 0.3320 - val_acc: 0.7358
Epoch 4/50
566006/566006 [=====] - 122s 216us/step - loss: 0.
3354 - acc: 0.7296 - val_loss: 0.3195 - val_acc: 0.7475
Epoch 5/50
566006/566006 [=====] - 123s 217us/step - loss: 0.
3230 - acc: 0.7413 - val_loss: 0.3086 - val_acc: 0.7552
Epoch 6/50
566006/566006 [=====] - 123s 217us/step - loss: 0.
3132 - acc: 0.7509 - val_loss: 0.3103 - val_acc: 0.7689
Epoch 7/50
566006/566006 [=====] - 123s 217us/step - loss: 0.
3048 - acc: 0.7592 - val_loss: 0.2966 - val_acc: 0.7700
Epoch 8/50
566006/566006 [=====] - 123s 217us/step - loss: 0.
2972 - acc: 0.7668 - val_loss: 0.2924 - val_acc: 0.7814
Epoch 9/50
566006/566006 [=====] - 123s 217us/step - loss: 0.
2908 - acc: 0.7733 - val_loss: 0.2932 - val_acc: 0.7839
Epoch 10/50
566006/566006 [=====] - 123s 217us/step - loss: 0.
2849 - acc: 0.7799 - val_loss: 0.2874 - val_acc: 0.7816
Epoch 11/50
566006/566006 [=====] - 123s 218us/step - loss: 0.
2796 - acc: 0.7844 - val_loss: 0.2858 - val_acc: 0.7958
Epoch 12/50
566006/566006 [=====] - 123s 218us/step - loss: 0.
2753 - acc: 0.7887 - val_loss: 0.2834 - val_acc: 0.7964
Epoch 13/50
566006/566006 [=====] - 123s 217us/step - loss: 0.
2715 - acc: 0.7928 - val_loss: 0.2860 - val_acc: 0.8042
Epoch 14/50
566006/566006 [=====] - 123s 217us/step - loss: 0.
2675 - acc: 0.7967 - val_loss: 0.2847 - val_acc: 0.8043
Epoch 15/50
566006/566006 [=====] - 123s 218us/step - loss: 0.
2638 - acc: 0.7998 - val_loss: 0.2817 - val_acc: 0.8035
Epoch 16/50
566006/566006 [=====] - 123s 218us/step - loss: 0.
2609 - acc: 0.8027 - val_loss: 0.2805 - val_acc: 0.8034
Epoch 17/50
566006/566006 [=====] - 123s 218us/step - loss: 0.
2573 - acc: 0.8062 - val_loss: 0.2834 - val_acc: 0.8087

```
Epoch 18/50
566006/566006 [=====] - 123s 218us/step - loss: 0.
2547 - acc: 0.8089 - val_loss: 0.2764 - val_acc: 0.8081
Epoch 19/50
566006/566006 [=====] - 123s 218us/step - loss: 0.
2513 - acc: 0.8117 - val_loss: 0.2773 - val_acc: 0.8107
Epoch 20/50
566006/566006 [=====] - 123s 218us/step - loss: 0.
2493 - acc: 0.8140 - val_loss: 0.2748 - val_acc: 0.8081
Epoch 21/50
566006/566006 [=====] - 123s 218us/step - loss: 0.
2471 - acc: 0.8158 - val_loss: 0.2771 - val_acc: 0.8129
Epoch 22/50
566006/566006 [=====] - 123s 218us/step - loss: 0.
2448 - acc: 0.8178 - val_loss: 0.2758 - val_acc: 0.8116
Epoch 23/50
566006/566006 [=====] - 123s 218us/step - loss: 0.
2436 - acc: 0.8189 - val_loss: 0.2772 - val_acc: 0.8122
Epoch 24/50
566006/566006 [=====] - 123s 218us/step - loss: 0.
2413 - acc: 0.8219 - val_loss: 0.2753 - val_acc: 0.8138
Epoch 25/50
566006/566006 [=====] - 124s 218us/step - loss: 0.
2395 - acc: 0.8230 - val_loss: 0.2841 - val_acc: 0.8189
Epoch 26/50
566006/566006 [=====] - 123s 218us/step - loss: 0.
2385 - acc: 0.8243 - val_loss: 0.2749 - val_acc: 0.8131
Epoch 27/50
566006/566006 [=====] - 123s 217us/step - loss: 0.
2365 - acc: 0.8259 - val_loss: 0.2791 - val_acc: 0.8185
Epoch 28/50
566006/566006 [=====] - 123s 218us/step - loss: 0.
2353 - acc: 0.8275 - val_loss: 0.2808 - val_acc: 0.8180
Epoch 29/50
566006/566006 [=====] - 123s 218us/step - loss: 0.
2334 - acc: 0.8287 - val_loss: 0.2702 - val_acc: 0.8133
Epoch 30/50
566006/566006 [=====] - 123s 218us/step - loss: 0.
2327 - acc: 0.8297 - val_loss: 0.2763 - val_acc: 0.8161
Epoch 31/50
566006/566006 [=====] - 124s 218us/step - loss: 0.
2319 - acc: 0.8304 - val_loss: 0.2752 - val_acc: 0.8168
Epoch 32/50
566006/566006 [=====] - 124s 218us/step - loss: 0.
2301 - acc: 0.8321 - val_loss: 0.2795 - val_acc: 0.8211
Epoch 33/50
566006/566006 [=====] - 124s 219us/step - loss: 0.
2295 - acc: 0.8326 - val_loss: 0.2750 - val_acc: 0.8210
Epoch 34/50
566006/566006 [=====] - 124s 218us/step - loss: 0.
2275 - acc: 0.8343 - val_loss: 0.2788 - val_acc: 0.8235
Epoch 35/50
566006/566006 [=====] - 124s 218us/step - loss: 0.
2266 - acc: 0.8351 - val_loss: 0.2787 - val_acc: 0.8218
Epoch 36/50
566006/566006 [=====] - 124s 218us/step - loss: 0.
2267 - acc: 0.8355 - val_loss: 0.2731 - val_acc: 0.8206
```

```
Epoch 37/50
566006/566006 [=====] - 124s 219us/step - loss: 0.
2250 - acc: 0.8363 - val_loss: 0.2828 - val_acc: 0.8224
Epoch 38/50
566006/566006 [=====] - 124s 219us/step - loss: 0.
2248 - acc: 0.8370 - val_loss: 0.2737 - val_acc: 0.8194
Epoch 39/50
566006/566006 [=====] - 124s 218us/step - loss: 0.
2244 - acc: 0.8371 - val_loss: 0.2847 - val_acc: 0.8268
Epoch 40/50
566006/566006 [=====] - 123s 218us/step - loss: 0.
2233 - acc: 0.8386 - val_loss: 0.2738 - val_acc: 0.8225
Epoch 41/50
566006/566006 [=====] - 124s 219us/step - loss: 0.
2218 - acc: 0.8402 - val_loss: 0.2813 - val_acc: 0.8244
Epoch 42/50
566006/566006 [=====] - 124s 219us/step - loss: 0.
2208 - acc: 0.8404 - val_loss: 0.2774 - val_acc: 0.8256
Epoch 43/50
566006/566006 [=====] - 124s 218us/step - loss: 0.
2203 - acc: 0.8409 - val_loss: 0.2760 - val_acc: 0.8248
Epoch 44/50
566006/566006 [=====] - 124s 219us/step - loss: 0.
2196 - acc: 0.8419 - val_loss: 0.2792 - val_acc: 0.8262
Epoch 45/50
566006/566006 [=====] - 124s 218us/step - loss: 0.
2192 - acc: 0.8417 - val_loss: 0.2807 - val_acc: 0.8257
Epoch 46/50
566006/566006 [=====] - 124s 219us/step - loss: 0.
2182 - acc: 0.8427 - val_loss: 0.2778 - val_acc: 0.8241
Epoch 47/50
566006/566006 [=====] - 124s 219us/step - loss: 0.
2176 - acc: 0.8437 - val_loss: 0.2774 - val_acc: 0.8257
Epoch 48/50
566006/566006 [=====] - 124s 219us/step - loss: 0.
2178 - acc: 0.8436 - val_loss: 0.2718 - val_acc: 0.8221
Epoch 49/50
566006/566006 [=====] - 124s 219us/step - loss: 0.
2169 - acc: 0.8450 - val_loss: 0.2786 - val_acc: 0.8265
Epoch 50/50
566006/566006 [=====] - 123s 218us/step - loss: 0.
2159 - acc: 0.8451 - val_loss: 0.2819 - val_acc: 0.8277
```

```
In [11]: def add_attention_after_lstm(inputs, SINGLE_ATTENTION_VECTOR=False):
    a = Lambda(lambda x: KB.expand_dims(x, axis=1))(inputs)
    input_dim = int(a.shape[2])
    a = Permute((2, 1))(a)
    a = Reshape((input_dim, 1))(a) # this line is not useful. It's just to know which dimension is what.
    a = Dense(1, activation='softmax')(a)
    if SINGLE_ATTENTION_VECTOR:
        a = Lambda(lambda x: KB.mean(x, axis=1))(a)
        a = RepeatVector(input_dim)(a)
    a_probs = Permute((2, 1))(a)
    a_probs = Lambda(lambda x: KB.squeeze(x, axis=1))(a_probs)
    outputs = KL.multiply([inputs, a_probs])
    return outputs
```

```
In [12]: def add_attention_before_lstm(inputs, T, SINGLE_ATTENTION_VECTOR=True):
    a = inputs
    input_dim = int(a.shape[2])
    a = Permute((2, 1))(a)
    a = Reshape((input_dim, T))(a) # this line is not useful. It's just to know which dimension is what.
    a = Dense(T, activation='softmax')(a)
    if SINGLE_ATTENTION_VECTOR:
        a = Lambda(lambda x: KB.mean(x, axis=1))(a)
        a = RepeatVector(input_dim)(a)
    a_probs = Permute((2, 1))(a)
    outputs = KL.multiply([inputs, a_probs])
    return outputs
```

```
In [13]: embedding_layer_att_before = Embedding(nb_words, EMBEDDING_DIM, weights=[embedding_matrix],
                                             input_length=MAX_SEQUENCE_LENGTH, trainable=False)
lstm_layer_att_before = LSTM(num_lstm, dropout=rate_drop_lstm, recurrent_dropout=rate_drop_lstm)
sequence_1_input_att_before = Input(shape=(MAX_SEQUENCE_LENGTH,), dtype='int32')
embedded_sequences_1_att_before = embedding_layer_att_before(sequence_1_input_att_before)
x1_att_before = add_attention_before_lstm(embedded_sequences_1_att_before, MAX_SEQUENCE_LENGTH, SINGLE_ATTENTION_VECTOR=True)
x1_att_before = lstm_layer_att_before(x1_att_before)
sequence_2_input_att_before = Input(shape=(MAX_SEQUENCE_LENGTH,), dtype='int32')
embedded_sequences_2_att_before = embedding_layer_att_before(sequence_2_input_att_before)
y1_att_before = add_attention_before_lstm(embedded_sequences_2_att_before, MAX_SEQUENCE_LENGTH, SINGLE_ATTENTION_VECTOR=True)
y1_att_before = lstm_layer_att_before(embedded_sequences_2_att_before)
merged_att_before = concatenate([x1_att_before, y1_att_before])
merged_att_before = Dropout(rate_drop_dense)(merged_att_before)
merged_att_before = BatchNormalization()(merged_att_before)
merged_att_before = Dense(num_dense, activation=act)(merged_att_before)
merged_att_before = Dropout(rate_drop_dense)(merged_att_before)
merged_att_before = BatchNormalization()(merged_att_before)
preds_att_before = Dense(1, activation='sigmoid')(merged_att_before)
```

```
In [14]: model_att_before = Model(inputs=[sequence_1_input_att_before, sequence_2_input_att_before], outputs=preds_att_before)
model_att_before.compile(loss='binary_crossentropy', optimizer='adam', metrics=['acc'])
model_att_before.summary()
bst_model_path_att_before = 'att_before_' + STAMP + '.h5'
model_checkpoint_att_before = ModelCheckpoint(bst_model_path_att_before, save_best_only=True, save_weights_only=True)
if LOAD_DATA:
    model_att_before.load_weights(bst_model_path_att_before)
hist_att_before = model_att_before.fit([data_1_train, data_2_train], labels_train, \
    validation_data=([data_1_val, data_2_val], labels_val, weight_val), \
    epochs=50, batch_size=2048, shuffle=True, \
    class_weight=class_weight, callbacks=[model_checkpoint_att_before])
model_att_before.load_weights(bst_model_path_att_before)
bst_val_score_att_before = min(hist_att_before.history['val_loss'])
```

Layer (type)	Output Shape	Param #	Connected to
input_3 (InputLayer)	(None, 30)	0	
embedding_2 (Embedding)	(None, 30, 300)	25655700	input_3[0]
[0]			input_4[0]
[0]			
permute_1 (Permute)	(None, 300, 30)	0	embedding_2[0][0]
[0]			
reshape_1 (Reshape)	(None, 300, 30)	0	permute_1[0]
[0]			
dense_3 (Dense)	(None, 300, 30)	930	reshape_1[0]
[0]			
lambda_1 (Lambda)	(None, 30)	0	dense_3[0]
[0]			
repeat_vector_1 (RepeatVector)	(None, 300, 30)	0	lambda_1[0]
[0]			
permute_2 (Permute)	(None, 30, 300)	0	repeat_vector_1[0]
[0]			
input_4 (InputLayer)	(None, 30)	0	
multiply_1 (Multiply)	(None, 30, 300)	0	embedding_2[0][0]
[0]			permute_2
[0]			
lstm_2 (LSTM)	(None, 265)	599960	multiply_1[0]
[0]			embedding_2[1]
[0]			
concatenate_2 (Concatenate)	(None, 530)	0	lstm_2[0]

[0]			lstm_2[1]
[0]			
dropout_3 (Dropout)	(None, 530)	0	concatenat
e_2[0][0]			
batch_normalization_3 (BatchNor	(None, 530)	2120	dropout_3
[0][0]			
dense_5 (Dense)	(None, 145)	76995	batch_norm
alization_3[0][0]			
dropout_4 (Dropout)	(None, 145)	0	dense_5[0]
[0]			
batch_normalization_4 (BatchNor	(None, 145)	580	dropout_4
[0][0]			
dense_6 (Dense)	(None, 1)	146	batch_norm
alization_4[0][0]			
<hr/>			
<hr/>			
Total params: 26,336,431			
Trainable params: 679,381			
Non-trainable params: 25,657,050			
<hr/>			
Train on 566006 samples, validate on 242574 samples			
Epoch 1/50			
566006/566006 [=====] - 151s 267us/step - loss: 0.			
4538 - acc: 0.6585 - val_loss: 0.3901 - val_acc: 0.6979			
Epoch 2/50			
566006/566006 [=====] - 173s 306us/step - loss: 0.			
3881 - acc: 0.6855 - val_loss: 0.3642 - val_acc: 0.6993			
Epoch 3/50			
566006/566006 [=====] - 184s 326us/step - loss: 0.			
3714 - acc: 0.6985 - val_loss: 0.3593 - val_acc: 0.7187			
Epoch 4/50			
566006/566006 [=====] - 172s 303us/step - loss: 0.			
3605 - acc: 0.7072 - val_loss: 0.3429 - val_acc: 0.7060			
Epoch 5/50			
566006/566006 [=====] - 148s 262us/step - loss: 0.			
3517 - acc: 0.7161 - val_loss: 0.3344 - val_acc: 0.7240			
Epoch 6/50			
566006/566006 [=====] - 149s 263us/step - loss: 0.			
3452 - acc: 0.7222 - val_loss: 0.3322 - val_acc: 0.7377			
Epoch 7/50			
566006/566006 [=====] - 149s 263us/step - loss: 0.			
3394 - acc: 0.7278 - val_loss: 0.3268 - val_acc: 0.7448			
Epoch 8/50			

```
566006/566006 [=====] - 149s 263us/step - loss: 0.  
3349 - acc: 0.7327 - val_loss: 0.3240 - val_acc: 0.7500  
Epoch 9/50  
566006/566006 [=====] - 149s 263us/step - loss: 0.  
3306 - acc: 0.7369 - val_loss: 0.3177 - val_acc: 0.7496  
Epoch 10/50  
566006/566006 [=====] - 149s 263us/step - loss: 0.  
3264 - acc: 0.7415 - val_loss: 0.3153 - val_acc: 0.7414  
Epoch 11/50  
566006/566006 [=====] - 148s 262us/step - loss: 0.  
3234 - acc: 0.7441 - val_loss: 0.3173 - val_acc: 0.7628  
Epoch 12/50  
566006/566006 [=====] - 148s 262us/step - loss: 0.  
3198 - acc: 0.7482 - val_loss: 0.3111 - val_acc: 0.7603  
Epoch 13/50  
566006/566006 [=====] - 148s 262us/step - loss: 0.  
3172 - acc: 0.7509 - val_loss: 0.3125 - val_acc: 0.7683  
Epoch 14/50  
566006/566006 [=====] - 148s 262us/step - loss: 0.  
3142 - acc: 0.7542 - val_loss: 0.3061 - val_acc: 0.7470  
Epoch 15/50  
566006/566006 [=====] - 148s 262us/step - loss: 0.  
3114 - acc: 0.7567 - val_loss: 0.3112 - val_acc: 0.7696  
Epoch 16/50  
566006/566006 [=====] - 149s 263us/step - loss: 0.  
3087 - acc: 0.7596 - val_loss: 0.3073 - val_acc: 0.7652  
Epoch 17/50  
566006/566006 [=====] - 149s 263us/step - loss: 0.  
3062 - acc: 0.7620 - val_loss: 0.3047 - val_acc: 0.7721  
Epoch 18/50  
566006/566006 [=====] - 149s 263us/step - loss: 0.  
3039 - acc: 0.7644 - val_loss: 0.3041 - val_acc: 0.7741  
Epoch 19/50  
566006/566006 [=====] - 149s 263us/step - loss: 0.  
3028 - acc: 0.7650 - val_loss: 0.3019 - val_acc: 0.7706  
Epoch 20/50  
566006/566006 [=====] - 149s 264us/step - loss: 0.  
3006 - acc: 0.7676 - val_loss: 0.3017 - val_acc: 0.7732  
Epoch 21/50  
566006/566006 [=====] - 149s 263us/step - loss: 0.  
2991 - acc: 0.7693 - val_loss: 0.3030 - val_acc: 0.7762  
Epoch 22/50  
566006/566006 [=====] - 149s 263us/step - loss: 0.  
2965 - acc: 0.7720 - val_loss: 0.2969 - val_acc: 0.7710  
Epoch 23/50  
566006/566006 [=====] - 148s 262us/step - loss: 0.  
2951 - acc: 0.7727 - val_loss: 0.2978 - val_acc: 0.7689  
Epoch 24/50  
566006/566006 [=====] - 148s 261us/step - loss: 0.  
2935 - acc: 0.7748 - val_loss: 0.2948 - val_acc: 0.7765  
Epoch 25/50  
566006/566006 [=====] - 148s 261us/step - loss: 0.  
2928 - acc: 0.7754 - val_loss: 0.2966 - val_acc: 0.7805  
Epoch 26/50  
566006/566006 [=====] - 149s 262us/step - loss: 0.  
2906 - acc: 0.7780 - val_loss: 0.2965 - val_acc: 0.7800  
Epoch 27/50
```

```
566006/566006 [=====] - 148s 262us/step - loss: 0.  
2893 - acc: 0.7794 - val_loss: 0.2970 - val_acc: 0.7876  
Epoch 28/50  
566006/566006 [=====] - 148s 261us/step - loss: 0.  
2880 - acc: 0.7805 - val_loss: 0.2914 - val_acc: 0.7789  
Epoch 29/50  
566006/566006 [=====] - 148s 262us/step - loss: 0.  
2871 - acc: 0.7819 - val_loss: 0.2941 - val_acc: 0.7836  
Epoch 30/50  
566006/566006 [=====] - 148s 262us/step - loss: 0.  
2854 - acc: 0.7831 - val_loss: 0.2952 - val_acc: 0.7853  
Epoch 31/50  
566006/566006 [=====] - 148s 262us/step - loss: 0.  
2847 - acc: 0.7834 - val_loss: 0.2944 - val_acc: 0.7897  
Epoch 32/50  
566006/566006 [=====] - 148s 262us/step - loss: 0.  
2837 - acc: 0.7847 - val_loss: 0.2917 - val_acc: 0.7826  
Epoch 33/50  
566006/566006 [=====] - 148s 261us/step - loss: 0.  
2823 - acc: 0.7860 - val_loss: 0.2927 - val_acc: 0.7870  
Epoch 34/50  
566006/566006 [=====] - 148s 262us/step - loss: 0.  
2814 - acc: 0.7874 - val_loss: 0.2921 - val_acc: 0.7827  
Epoch 35/50  
566006/566006 [=====] - 148s 262us/step - loss: 0.  
2798 - acc: 0.7882 - val_loss: 0.2903 - val_acc: 0.7873  
Epoch 36/50  
566006/566006 [=====] - 148s 262us/step - loss: 0.  
2798 - acc: 0.7889 - val_loss: 0.2895 - val_acc: 0.7833  
Epoch 37/50  
566006/566006 [=====] - 148s 261us/step - loss: 0.  
2786 - acc: 0.7901 - val_loss: 0.2891 - val_acc: 0.7873  
Epoch 38/50  
566006/566006 [=====] - 148s 261us/step - loss: 0.  
2775 - acc: 0.7911 - val_loss: 0.2899 - val_acc: 0.7897  
Epoch 39/50  
566006/566006 [=====] - 148s 262us/step - loss: 0.  
2765 - acc: 0.7918 - val_loss: 0.2878 - val_acc: 0.7871  
Epoch 40/50  
566006/566006 [=====] - 148s 261us/step - loss: 0.  
2755 - acc: 0.7928 - val_loss: 0.2884 - val_acc: 0.7892  
Epoch 41/50  
566006/566006 [=====] - 148s 262us/step - loss: 0.  
2753 - acc: 0.7927 - val_loss: 0.2868 - val_acc: 0.7906  
Epoch 42/50  
566006/566006 [=====] - 148s 261us/step - loss: 0.  
2738 - acc: 0.7941 - val_loss: 0.2894 - val_acc: 0.7953  
Epoch 43/50  
566006/566006 [=====] - 149s 262us/step - loss: 0.  
2740 - acc: 0.7942 - val_loss: 0.2920 - val_acc: 0.7961  
Epoch 44/50  
566006/566006 [=====] - 149s 262us/step - loss: 0.  
2730 - acc: 0.7954 - val_loss: 0.2844 - val_acc: 0.7860  
Epoch 45/50  
566006/566006 [=====] - 148s 262us/step - loss: 0.  
2722 - acc: 0.7961 - val_loss: 0.2910 - val_acc: 0.7963  
Epoch 46/50
```

```
566006/566006 [=====] - 149s 263us/step - loss: 0.
2715 - acc: 0.7964 - val_loss: 0.2880 - val_acc: 0.7951
Epoch 47/50
566006/566006 [=====] - 148s 261us/step - loss: 0.
2707 - acc: 0.7974 - val_loss: 0.2894 - val_acc: 0.7981
Epoch 48/50
566006/566006 [=====] - 148s 262us/step - loss: 0.
2704 - acc: 0.7977 - val_loss: 0.2869 - val_acc: 0.7960
Epoch 49/50
566006/566006 [=====] - 148s 262us/step - loss: 0.
2695 - acc: 0.7989 - val_loss: 0.2838 - val_acc: 0.7889
Epoch 50/50
566006/566006 [=====] - 148s 262us/step - loss: 0.
2689 - acc: 0.7987 - val_loss: 0.2878 - val_acc: 0.8000
```

```
In [15]: embedding_layer_att_after = Embedding(nb_words, EMBEDDING_DIM, weights=[embedding_matrix],
                                             input_length=MAX_SEQUENCE_LENGTH, trainable=False)
lstm_layer_att_after = LSTM(num_lstm, dropout=rate_drop_lstm, recurrent_dropout=rate_drop_lstm)
sequence_1_input_att_after = Input(shape=(MAX_SEQUENCE_LENGTH,), dtype='int32')
embedded_sequences_1_att_after = embedding_layer_att_after(sequence_1_input_att_after)
x1_att_after = lstm_layer_att_after(embedded_sequences_1_att_after)
x1_att_after = add_attention_after_lstm(x1_att_after, SINGLE_ATTENTION_VECTOR=False)
sequence_2_input_att_after = Input(shape=(MAX_SEQUENCE_LENGTH,), dtype='int32')
embedded_sequences_2_att_after = embedding_layer_att_after(sequence_2_input_att_after)
y1_att_after = lstm_layer_att_after(embedded_sequences_2_att_after)
y1_att_after = add_attention_after_lstm(y1_att_after, SINGLE_ATTENTION_VECTOR=False)
merged_att_after = concatenate([x1_att_after, y1_att_after])
merged_att_after = Dropout(rate_drop_dense)(merged_att_after)
merged_att_after = BatchNormalization()(merged_att_after)
merged_att_after = Dense(num_dense, activation=act)(merged_att_after)
merged_att_after = Dropout(rate_drop_dense)(merged_att_after)
merged_att_after = BatchNormalization()(merged_att_after)
preds_att_after = Dense(1, activation='sigmoid')(merged_att_after)
```

```
In [16]: model_att_after = Model(inputs=[sequence_1_input_att_after, sequence_2_input_att_after], outputs=preds_att_after)
model_att_after.compile(loss='binary_crossentropy', optimizer='adam', metrics=['acc'])
model_att_after.summary()
bst_model_path_att_after = 'att_after_' + STAMP + '.h5'
model_checkpoint_att_after = ModelCheckpoint(bst_model_path_att_after, save_best_only=True, save_weights_only=True)
if LOAD_DATA:
    model_att_after.load_weights(bst_model_path_att_after)
hist_att_after = model_att_after.fit([data_1_train, data_2_train], labels_train, \
    validation_data=([data_1_val, data_2_val], labels_val, weight_val), \
    epochs=50, batch_size=2048, shuffle=True, \
    class_weight=class_weight, callbacks=[model_checkpoint_att_after])
model_att_after.load_weights(bst_model_path_att_after)
bst_val_score_att_after = min(hist_att_after.history['val_loss'])
```

Layer (type) to	Output Shape	Param #	Connected
input_5 (InputLayer)	(None, 30)	0	
input_6 (InputLayer)	(None, 30)	0	
embedding_3 (Embedding) [0]	(None, 30, 300)	25655700	input_5[0] input_6[0]
lstm_3 (LSTM) 3[0][0]	(None, 265)	599960	embedding_3[1][0]
lambda_3 (Lambda) [0]	(None, 1, 265)	0	lstm_3[0]
lambda_5 (Lambda) [0]	(None, 1, 265)	0	lstm_3[1]
permute_5 (Permute) [0][0]	(None, 265, 1)	0	lambda_3
permute_7 (Permute) [0][0]	(None, 265, 1)	0	lambda_5
reshape_3 (Reshape) [0][0]	(None, 265, 1)	0	permute_5
reshape_4 (Reshape) [0][0]	(None, 265, 1)	0	permute_7
dense_7 (Dense) [0][0]	(None, 265, 1)	2	reshape_3
dense_8 (Dense) [0][0]	(None, 265, 1)	2	reshape_4

permute_6 (Permute) [0]	(None, 1, 265)	0	dense_7[0]
permute_8 (Permute) [0]	(None, 1, 265)	0	dense_8[0]
lambda_4 (Lambda) [0][0]	(None, 265)	0	permute_6
lambda_6 (Lambda) [0][0]	(None, 265)	0	permute_8
multiply_3 (Multiply) [0] [0][0]	(None, 265)	0	lstm_3[0] lambda_4
multiply_4 (Multiply) [0] [0][0]	(None, 265)	0	lstm_3[1] lambda_6
concatenate_3 (Concatenate) [0][0] [0][0]	(None, 530)	0	multiply_3 multiply_4
dropout_5 (Dropout) e_3[0][0]	(None, 530)	0	concatenat
batch_normalization_5 (BatchNor [0][0]	(None, 530)	2120	dropout_5
dense_9 (Dense) alization_5[0][0]	(None, 145)	76995	batch_norm
dropout_6 (Dropout) [0]	(None, 145)	0	dense_9[0]
batch_normalization_6 (BatchNor [0][0]	(None, 145)	580	dropout_6
dense_10 (Dense) alization_6[0][0]	(None, 1)	146	batch_norm

```
=====
Total params: 26,335,505
Trainable params: 678,455
Non-trainable params: 25,657,050
```

```
Train on 566006 samples, validate on 242574 samples
Epoch 1/50
566006/566006 [=====] - 128s 227us/step - loss: 0.
4397 - acc: 0.6695 - val_loss: 0.3685 - val_acc: 0.7189
Epoch 2/50
566006/566006 [=====] - 126s 222us/step - loss: 0.
3706 - acc: 0.6986 - val_loss: 0.3455 - val_acc: 0.7106
Epoch 3/50
566006/566006 [=====] - 125s 222us/step - loss: 0.
3497 - acc: 0.7161 - val_loss: 0.3385 - val_acc: 0.7361
Epoch 4/50
566006/566006 [=====] - 125s 222us/step - loss: 0.
3340 - acc: 0.7302 - val_loss: 0.3213 - val_acc: 0.7507
Epoch 5/50
566006/566006 [=====] - 126s 222us/step - loss: 0.
3218 - acc: 0.7420 - val_loss: 0.3129 - val_acc: 0.7648
Epoch 6/50
566006/566006 [=====] - 126s 222us/step - loss: 0.
3126 - acc: 0.7511 - val_loss: 0.3013 - val_acc: 0.7669
Epoch 7/50
566006/566006 [=====] - 126s 222us/step - loss: 0.
3040 - acc: 0.7600 - val_loss: 0.2993 - val_acc: 0.7708
Epoch 8/50
566006/566006 [=====] - 126s 222us/step - loss: 0.
2971 - acc: 0.7661 - val_loss: 0.2948 - val_acc: 0.7887
Epoch 9/50
566006/566006 [=====] - 126s 222us/step - loss: 0.
2906 - acc: 0.7727 - val_loss: 0.2865 - val_acc: 0.7813
Epoch 10/50
566006/566006 [=====] - 126s 222us/step - loss: 0.
2847 - acc: 0.7791 - val_loss: 0.2880 - val_acc: 0.7905
Epoch 11/50
566006/566006 [=====] - 126s 222us/step - loss: 0.
2795 - acc: 0.7840 - val_loss: 0.2888 - val_acc: 0.8002
Epoch 12/50
566006/566006 [=====] - 126s 222us/step - loss: 0.
2747 - acc: 0.7886 - val_loss: 0.2806 - val_acc: 0.7989
Epoch 13/50
566006/566006 [=====] - 125s 222us/step - loss: 0.
2704 - acc: 0.7932 - val_loss: 0.2813 - val_acc: 0.8016
Epoch 14/50
566006/566006 [=====] - 126s 222us/step - loss: 0.
2665 - acc: 0.7970 - val_loss: 0.2802 - val_acc: 0.8041
Epoch 15/50
566006/566006 [=====] - 126s 222us/step - loss: 0.
2628 - acc: 0.8008 - val_loss: 0.2850 - val_acc: 0.8078
Epoch 16/50
566006/566006 [=====] - 126s 222us/step - loss: 0.
2593 - acc: 0.8041 - val_loss: 0.2782 - val_acc: 0.8044
Epoch 17/50
```

```
566006/566006 [=====] - 126s 222us/step - loss: 0.  
2562 - acc: 0.8069 - val_loss: 0.2799 - val_acc: 0.8113  
Epoch 18/50  
566006/566006 [=====] - 126s 222us/step - loss: 0.  
2535 - acc: 0.8093 - val_loss: 0.2770 - val_acc: 0.8024  
Epoch 19/50  
566006/566006 [=====] - 125s 222us/step - loss: 0.  
2506 - acc: 0.8120 - val_loss: 0.2770 - val_acc: 0.8114  
Epoch 20/50  
566006/566006 [=====] - 126s 222us/step - loss: 0.  
2482 - acc: 0.8146 - val_loss: 0.2805 - val_acc: 0.8153  
Epoch 21/50  
566006/566006 [=====] - 126s 222us/step - loss: 0.  
2467 - acc: 0.8157 - val_loss: 0.2833 - val_acc: 0.8154  
Epoch 22/50  
566006/566006 [=====] - 126s 222us/step - loss: 0.  
2442 - acc: 0.8183 - val_loss: 0.2794 - val_acc: 0.8173  
Epoch 23/50  
566006/566006 [=====] - 125s 221us/step - loss: 0.  
2428 - acc: 0.8200 - val_loss: 0.2805 - val_acc: 0.8167  
Epoch 24/50  
566006/566006 [=====] - 124s 220us/step - loss: 0.  
2406 - acc: 0.8215 - val_loss: 0.2759 - val_acc: 0.8159  
Epoch 25/50  
566006/566006 [=====] - 125s 221us/step - loss: 0.  
2393 - acc: 0.8234 - val_loss: 0.2812 - val_acc: 0.8205  
Epoch 26/50  
566006/566006 [=====] - 142s 250us/step - loss: 0.  
2371 - acc: 0.8248 - val_loss: 0.2772 - val_acc: 0.8166  
Epoch 27/50  
566006/566006 [=====] - 188s 333us/step - loss: 0.  
2356 - acc: 0.8267 - val_loss: 0.2801 - val_acc: 0.8193  
Epoch 28/50  
566006/566006 [=====] - 151s 267us/step - loss: 0.  
2343 - acc: 0.8279 - val_loss: 0.2768 - val_acc: 0.8182  
Epoch 29/50  
566006/566006 [=====] - 153s 270us/step - loss: 0.  
2332 - acc: 0.8287 - val_loss: 0.2738 - val_acc: 0.8187  
Epoch 30/50  
566006/566006 [=====] - 153s 271us/step - loss: 0.  
2321 - acc: 0.8297 - val_loss: 0.2843 - val_acc: 0.8248  
Epoch 31/50  
566006/566006 [=====] - 153s 270us/step - loss: 0.  
2301 - acc: 0.8314 - val_loss: 0.2779 - val_acc: 0.8214  
Epoch 32/50  
566006/566006 [=====] - 155s 273us/step - loss: 0.  
2294 - acc: 0.8323 - val_loss: 0.2760 - val_acc: 0.8206  
Epoch 33/50  
566006/566006 [=====] - 153s 270us/step - loss: 0.  
2289 - acc: 0.8330 - val_loss: 0.2769 - val_acc: 0.8220  
Epoch 34/50  
566006/566006 [=====] - 150s 265us/step - loss: 0.  
2270 - acc: 0.8343 - val_loss: 0.2778 - val_acc: 0.8226  
Epoch 35/50  
566006/566006 [=====] - 152s 269us/step - loss: 0.  
2259 - acc: 0.8354 - val_loss: 0.2751 - val_acc: 0.8231  
Epoch 36/50
```

```
566006/566006 [=====] - 157s 277us/step - loss: 0.  
2253 - acc: 0.8364 - val_loss: 0.2775 - val_acc: 0.8244  
Epoch 37/50  
566006/566006 [=====] - 154s 272us/step - loss: 0.  
2242 - acc: 0.8371 - val_loss: 0.2737 - val_acc: 0.8224  
Epoch 38/50  
566006/566006 [=====] - 153s 270us/step - loss: 0.  
2240 - acc: 0.8376 - val_loss: 0.2779 - val_acc: 0.8217  
Epoch 39/50  
566006/566006 [=====] - 151s 267us/step - loss: 0.  
2232 - acc: 0.8379 - val_loss: 0.2809 - val_acc: 0.8240  
Epoch 40/50  
566006/566006 [=====] - 157s 277us/step - loss: 0.  
2217 - acc: 0.8393 - val_loss: 0.2823 - val_acc: 0.8278  
Epoch 41/50  
566006/566006 [=====] - 130s 229us/step - loss: 0.  
2214 - acc: 0.8399 - val_loss: 0.2799 - val_acc: 0.8243  
Epoch 42/50  
566006/566006 [=====] - 125s 221us/step - loss: 0.  
2204 - acc: 0.8410 - val_loss: 0.2841 - val_acc: 0.8271  
Epoch 43/50  
566006/566006 [=====] - 125s 221us/step - loss: 0.  
2195 - acc: 0.8415 - val_loss: 0.2750 - val_acc: 0.8242  
Epoch 44/50  
566006/566006 [=====] - 126s 222us/step - loss: 0.  
2194 - acc: 0.8415 - val_loss: 0.2755 - val_acc: 0.8260  
Epoch 45/50  
566006/566006 [=====] - 126s 222us/step - loss: 0.  
2184 - acc: 0.8427 - val_loss: 0.2771 - val_acc: 0.8239  
Epoch 46/50  
566006/566006 [=====] - 126s 222us/step - loss: 0.  
2176 - acc: 0.8434 - val_loss: 0.2789 - val_acc: 0.8277  
Epoch 47/50  
566006/566006 [=====] - 126s 222us/step - loss: 0.  
2177 - acc: 0.8430 - val_loss: 0.2748 - val_acc: 0.8272  
Epoch 48/50  
566006/566006 [=====] - 126s 222us/step - loss: 0.  
2167 - acc: 0.8440 - val_loss: 0.2734 - val_acc: 0.8253  
Epoch 49/50  
566006/566006 [=====] - 125s 222us/step - loss: 0.  
2156 - acc: 0.8448 - val_loss: 0.2840 - val_acc: 0.8311  
Epoch 50/50  
566006/566006 [=====] - 126s 222us/step - loss: 0.  
2161 - acc: 0.8448 - val_loss: 0.2822 - val_acc: 0.8280
```